


Sparkrun intelligence scores (grey = test failed or could not run)




Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	official-qwen3.8-27b-fp8-vllm				
GSM8K strict	math reasoning	official-qwen3.8-27b-fp8-vllm				
GSM8K flex	math reasoning	official-qwen3.8-27b-fp8-vllm				
ARC acc	science reasoning	official-qwen3.8-27b-fp8-vllm				
ARC norm	science reasoning	official-qwen3.8-27b-fp8-vllm				
HellaSwag acc	commonsense	official-qwen3.8-27b-fp8-vllm				
HellaSwag norm	commonsense	official-qwen3.8-27b-fp8-vllm				
HumanEval	coding	official-qwen3.8-27b-fp8-vllm				
MBPP	coding	official-qwen3.8-27b-fp8-vllm				
BFCL overall	tool calling	official-qwen3.8-27b-fp8-vllm				

- Ranking rule:
- Failures counted as 0 (shown grey with a code)
 - FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
 - Top 5 chosen by score for each test
 - Ties broken by Mean, then current order
 - Mean covers lm-eval metrics only (excludes BFCL)


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size

