


Sparkrun intelligence scores, failures as zero




Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm		
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm		
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm		
ARC acc	science reasoning	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		
ARC norm	science reasoning	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		
HellaSwag acc	commonsense	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		
HellaSwag norm	commonsense	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		
HumanEval	coding	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		
MBPP	coding	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b		

Ranking rule:
- Failures counted as 0
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order

Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size



Sparkrun benchmark TPS by prompt size

