


Sparkrun intelligence scores (grey = test failed or could not run)

Benchmark

eugr-qwen3-coder-next-int4-autoround	0.84	0.92	0.92	0.60	0.64	0.54	0.71	0.71	0.77	0.27	0.74
nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	0.84	0.92	0.92	0.61	0.64	0.54	0.73	0.46	0.78	0.23	0.72
official-qwen3.6-27b-fp8-vllm	0.86	0.58	0.60	0.60	0.60	0.61	0.78	0.85	0.60	0.28	0.68
community-north-mini-code-1.0-fp8-vllm-XanuNetworks	0.79	0.91	0.91	0.56	0.59	0.49	0.64	0.56	0.62	0.28	0.67
community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	0.78	0.88	0.88	0.54	0.57	0.50	0.65	0.58	0.59	0.27	0.66
community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	0.87	0.58	0.62	0.60	0.61	0.61	0.76	0.62	0.50	0.10	0.64
community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	0.86	0.41	0.46	0.55	0.55	0.57	0.73	0.59	0.72	0.27	0.60
official-qwen3.6-35b-a3b-fp8-vllm	0.85	0.35	0.37	0.54	0.55	0.60	0.76	0.59	0.67	0.27	0.59
community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	0.84	0.37	0.38	0.47	0.48	0.58	0.75	0.51	0.60	0.28	0.55
eugr-nemotron-3-nano-nvfp4	0.71	0.82	0.47	0.45	0.48	0.52	0.66	0.22	0.60	0.27	0.55
aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	0.83	0.10	0.14	0.52	0.53	0.58	0.77	0.62	0.59	0.27	0.52
poolside-Laguna-S-2-1-NVFP4	0.55	0.86	0.52	0.30	0.33	0.38	0.52	0.41	0.73	0.21	0.51
aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	0.79	0.07	0.09	0.48	0.49	0.53	0.64	0.76	0.66	0.27	0.50
aeon-7-ornith-1.0-35b-uncensored-nvfp4-dflash-vllm	0.84	0.05	0.06	0.54	0.56	0.58	0.68	0.52	0.61	0.27	0.49
official-qwen3.6-27b-fp8-mtp-vllm	ERROR	0.58	0.60	0.26	0.29	0.26	0.27	0.71	0.40	0.28	0.37
official-qwen3.6-35b-a3b-fp8-mtp-vllm	0.28	0.35	0.37	0.20	0.25	0.22	0.32	0.51	0.50	0.27	0.33
atlas-qwen3.6-35b-a3b-fp8-nvfp4head	UNSUP	0.90	0.91	UNSUP	UNSUP	UNSUP	UNSUP	0.02	0.17	0.14	0.22
atlas-qwen3.6-35b-a3b-fp8-mtp	UNSUP	0.90	0.91	UNSUP	UNSUP	UNSUP	UNSUP	0.02	0.17	0.14	0.22
atlas-qwen3.5-35b-a3b-nvfp4	UNSUP	0.90	0.86	UNSUP	UNSUP	UNSUP	UNSUP	0.04	0.00	0.17	0.20
atlas-qwen3-vl-30b-a3b-nvfp4	UNSUP	0.11	0.10	UNSUP	UNSUP	UNSUP	UNSUP	0.02	0.00	0.02	0.03
atlas-qwen3.6-27b-nvfp4-prefill-record	UNSUP	HANG	HANG	UNSUP	UNSUP	UNSUP	UNSUP	0.02	0.11	0.14	0.01
atlas-mistral-small-4-119b-nvfp4	UNSUP	ERROR	ERROR	UNSUP	UNSUP	UNSUP	UNSUP	0.00	0.02	0.22	0.00
atlas-qwen3.5-27b-dense-nvfp4	UNSUP	ERROR	ERROR	UNSUP	UNSUP	UNSUP	UNSUP	0.01	ERROR	0.16	0.00
atlas-nemotron-3-nano-30b-a3b-nvfp4	UNSUP	0.00	0.00	UNSUP	UNSUP	UNSUP	UNSUP	0.00	0.00	0.09	0.00
atlas-qwen3-next-80b-a3b-nvfp4	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
atlas-nemotron-3-super-120b-a12b-nvfp4-batch4	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
atlas-qwen3.6-27b-fp8	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
eugr-openai-gpt-oss-120b	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
eugr-glm-4.7-flash-awq	UNSUP	CRASH	CRASH	UNSUP	UNSUP	UNSUP	UNSUP	CRASH	CRASH	0.10	0.00
atlas-qwen3.5-0.8b-bf16-atlas	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
atlas-qwen3.5-122b-a10b-nvfp4-single	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
community-qwen3.5-122b-a10b-int4fp8-mtp-vllm-Taremimi7710	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
colibri-glm-5.2-int4-mtp	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
deepreinforce-ai-ornith-1.0-35b-bf16-mtp-vllm	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	FAILED	0.00
	MMLU	GSM8K strict	GSM8K flex	ARC acc	ARC norm	HellaSwag acc	HellaSwag norm	HumanEval	MBPP	BFCL overall	Mean




Score

Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	atlas-qwen3.6-35b-a3b-fp8-nvfp4	atlas-qwen3.6-35b-a3b-fp8-mtp
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	atlas-qwen3.6-35b-a3b-fp8-mtp	atlas-qwen3.6-35b-a3b-fp8-nvfp4	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC acc	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC norm	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	official-qwen3.6-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeee
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	poolside-Laguna-S-2-1-NVFP4	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	eugr-nemotron-3-nano-nvfp4

- Ranking rule:
- Failures counted as 0 (shown grey with a code)
 - FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
 - Top 5 chosen by score for each test
 - Ties broken by Mean, then current order
 - Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark TPOT by prompt size

