


Sparkrun intelligence scores (grey = test failed or could not run)




Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	official-qwen3-coder-next-int4-autoround-vllm	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm
GSM8K strict	math reasoning	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.8-27b-fp8-mtp-vllm
GSM8K flex	math reasoning	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.8-27b-fp8-mtp-vllm
ARC acc	science reasoning	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-deepseek-v4-flash-0731-b1-vllm	eugr-deepseek-v4-flash-0731	official-qwen3-coder-next-int4-autoround-vllm	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm
ARC norm	science reasoning	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	official-qwen3-coder-next-int4-autoround-vllm	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm
HellaSwag acc	commonsense	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm
HellaSwag norm	commonsense	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm
HumanEval	coding	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm	official-qwen3.8-27b-fp8-mtp-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm
MBPP	coding	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.8-27b-fp8-mtp-vllm
BFCL 17 of 22	tool calling	official-qwen3.8-27b-fp8-mtp-vllm	official-deepseek-v4-flash-0731-b1-vllm	eugr-deepseek-v4-flash-0731	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	miaai-qwen3.8-flash-next-nvfp4-mtp3-2x-vllm

- Ranking rule:
- Failures counted as 0 (shown grey with a code)
 - FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
 - Top 5 chosen by score for each test
 - Ties broken by Mean, then current order
 - Mean covers lm-eval metrics only (excludes BFCL)


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size

