


Sparkrun intelligence scores, failures as zero

Benchmark

eugr-qwen3-coder-next-int4-autoround	0.84	0.92	0.92	0.60	0.64	0.54	0.71	0.71	0.77	0.27	0.74
nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	0.84	0.92	0.92	0.61	0.64	0.54	0.73	0.46	0.78	0.23	0.72
official-qwen3.6-27b-fp8-vllm	0.86	0.58	0.60	0.60	0.60	0.61	0.78	0.85	0.60	0.28	0.68
community-north-mini-code-1.0-fp8-vllm-XanuNetworks	0.79	0.91	0.91	0.56	0.59	0.49	0.64	0.56	0.62	0.28	0.67
community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	0.78	0.88	0.88	0.54	0.57	0.50	0.65	0.58	0.59	0.27	0.66
community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	0.87	0.58	0.62	0.60	0.61	0.61	0.76	0.62	0.50	0.10	0.64
community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	0.86	0.41	0.46	0.55	0.55	0.57	0.73	0.59	0.72	0.27	0.60
official-qwen3.6-35b-a3b-fp8-vllm	0.85	0.35	0.37	0.54	0.55	0.60	0.76	0.59	0.67	0.27	0.59
community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	0.84	0.37	0.38	0.47	0.48	0.58	0.75	0.51	0.60	0.28	0.55
eugr-nemotron-3-nano-nvfp4	0.71	0.82	0.47	0.45	0.48	0.52	0.66	0.22	0.60	0.27	0.55
aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	0.83	0.10	0.14	0.52	0.53	0.58	0.77	0.62	0.59	0.27	0.52
aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	0.79	0.07	0.09	0.48	0.49	0.53	0.64	0.76	0.66	0.27	0.50
aeon-7-ornith-1.0-35b-uncensored-nvfp4-dflash-vllm	0.84	0.05	0.06	0.54	0.56	0.58	0.68	0.52	0.61	0.27	0.49
official-qwen3.6-27b-fp8-mtp-vllm	0.00	0.58	0.60	0.26	0.29	0.26	0.27	0.71	0.40	0.28	0.37
official-qwen3.6-35b-a3b-fp8-mtp-vllm	0.28	0.35	0.37	0.20	0.25	0.22	0.32	0.51	0.50	0.27	0.33
atlas-qwen3.6-35b-a3b-fp8-nvfp4head	0.00	0.90	0.91	0.00	0.00	0.00	0.00	0.02	0.17	0.14	0.22
atlas-qwen3.6-35b-a3b-fp8-mtp	0.00	0.90	0.91	0.00	0.00	0.00	0.00	0.02	0.17	0.14	0.22
atlas-qwen3.5-35b-a3b-nvfp4	0.00	0.90	0.86	0.00	0.00	0.00	0.00	0.04	0.00	0.17	0.20
atlas-qwen3-vl-30b-a3b-nvfp4	0.00	0.11	0.10	0.00	0.00	0.00	0.00	0.02	0.00	0.02	0.03
atlas-qwen3.6-27b-nvfp4-prefill-record	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.02	0.11	0.14	0.01
atlas-mistral-small-4-119b-nvfp4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.02	0.22	0.00
atlas-qwen3.5-27b-dense-nvfp4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.01	0.00	0.16	0.00
atlas-nemotron-3-nano-30b-a3b-nvfp4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.09	0.00
eugr-openai-gpt-oss-120b	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
atlas-nemotron-3-super-120b-a12b-nvfp4-batch4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
deepreinforce-ai-ornith-1.0-35b-bf16-mtp-vllm	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
atlas-qwen3-next-80b-a3b-nvfp4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-glm-4.7-flash-awq	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.10	0.00
atlas-qwen3.5-0.8b-bf16-atlas	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
atlas-qwen3.5-122b-a10b-nvfp4-single	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
community-qwen3.5-122b-a10b-int4fp8-mtp-vllm-Taremimi7710	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
atlas-qwen3.6-27b-fp8	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
colibri-glm-5.2-int4-mtp	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
	MMLU	GSM8K strict	GSM8K flex	ARC acc	ARC norm	HellaSwag acc	HellaSwag norm	HumanEval	MBPP	BFCL overall	Mean




Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	atlas-qwen3.6-35b-a3b-fp8-nvfp4	atlas-qwen3.6-35b-a3b-fp8-mtp
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	atlas-qwen3.6-35b-a3b-fp8-mtp	atlas-qwen3.6-35b-a3b-fp8-nvfp4	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC acc	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC norm	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	official-qwen3.6-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-baeeee
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	eugr-nemotron-3-nano-nvfp4

- Ranking rule:
- Failures counted as 0
 - Top 5 chosen by score for each test
 - Ties broken by Mean, then current order
 - Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



The chart displays the relationship between the number of prompt tokens and the time taken by various models to process them. The x-axis represents the number of prompt tokens, ranging from 256 to 259,267. The y-axis represents the time taken in milliseconds (TPOT ms), ranging from 0 to 350. The models are represented by different colored lines with markers at each data point.

Key observations from the chart:

- The orange line (topmost) shows a significant increase in time as the number of tokens increases, reaching approximately 340 ms at 259,267 tokens.
- The red line (second from top) also shows a sharp increase, reaching approximately 300 ms at 129,506 tokens.
- The green line (third from top) shows a moderate increase, reaching approximately 80 ms at 259,267 tokens.
- The blue line (fourth from top) shows a slight increase, reaching approximately 75 ms at 259,267 tokens.
- The purple line (fifth from top) shows a slight increase, reaching approximately 90 ms at 259,267 tokens.
- The brown line (sixth from top) shows a slight increase, reaching approximately 30 ms at 259,267 tokens.
- The grey line (bottommost) shows a slight increase, reaching approximately 25 ms at 259,267 tokens.

Prompt tokens	Orange	Red	Green	Blue	Purple	Brown	Grey
256	130	45	25	55	20	20	15
1,024	130	50	20	55	20	20	15
4,096	130	45	20	55	20	20	15
7,680	130	60	20	55	20	20	15
16,384	130	110	20	55	20	20	15
32,185	135	180	20	55	20	20	15
36,240	140	240	20	55	20	20	15
65,536	150	290	20	55	20	20	15
129,506	250	300	20	55	20	20	15
259,267	340	300	80	75	90	30	25

Prompt tokens