


Sparkrun intelligence scores (grey = test failed or could not run)

Benchmark




Score

Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	poolside-Laguna-S-2-1-NVFP4
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	eugr-nemotron-3.5-lightning
ARC acc	science reasoning	eugr-nemotron-3.5-lightning	nvidia-nemotron-3-puzzle-75b-a9b	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
ARC norm	science reasoning	eugr-nemotron-3.5-lightning	nvidia-nemotron-3-puzzle-75b-a9b	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncenvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b	eugr-qwen3-coder-next-int4-autoround	poolside-Laguna-S-2-1-NVFP4	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	eugr-nemotron-3.5-lightning
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	eugr-nemotron-3.5-lightning	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui

- Ranking rule:
- Failures counted as 0 (shown grey with a code)
 - FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
 - Top 5 chosen by score for each test
 - Ties broken by Mean, then current order
 - Mean covers lm-eval metrics only (excludes BFCL)


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



The chart displays the relationship between the number of prompt tokens and a specific metric for various models. The x-axis, labeled 'Prompt tokens', includes categories: 256B, 1K, 4K, 16K, 64K, 127K, and 254K. The y-axis ranges from 0 to 350. The data is represented by multiple colored lines, each corresponding to a different model. Most models show a relatively flat trend until 16K tokens, after which they exhibit a sharp increase. The orange line shows the most dramatic increase, reaching a value of approximately 340 at 254K tokens. The light green line also shows a significant increase, reaching about 305 at 127K tokens. Other models, such as the dark blue and purple lines, show more moderate increases, while several others remain relatively flat or show a slight upward trend.

Prompt tokens	Orange	Light Green	Dark Blue	Purple	Teal	Dark Green	Grey	Bright Green	Red	Pink	Light Blue	Dark Purple	Light Purple	Dark Red
256B	25	45	55	130	65	55	45	30	15	10	20	25	20	15
1K	25	55	55	130	65	55	50	30	15	10	20	25	20	15
4K	30	45	55	130	65	55	45	30	15	10	20	25	20	15
16K	80	110	30	130	65	55	50	30	15	10	20	25	20	15
64K	155	245	45	140	65	55	55	30	20	10	25	30	25	20
127K	250	305	70	155	70	55	60	40	30	10	25	30	35	25
254K	340	-	95	165	80	55	65	45	45	10	25	30	40	30

official-qwen3.6-27b-fp8-vllm-

official-qwen3.6-27b-fp8-mtp-vllm

nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm

community-qwen3.5-27b-fp8-dflash-vllm-banana-bae

poolside-Laguna-S-2-1-NVFP4
de-1_0-fp8-vllm-YanpuNetworks

aeon-7-gwen3.6-27b-uncensored-nvfp4-dflash-vllm

official-qwen3.6-35b-a3b-fp8-vllm

official-qwen3.6-35b-a3b-fp8-mtp-vllm
ever-nemotron-2-rp8-pyfp4

unity-north-mini-code-1.0-pyfp4-vllm-XanunNetworks

n3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie

eugr-qwen3-coder-next-int4-autoround-

community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
ougr-nemotron-3.5-lightning

aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm

aeon-7-ornith-1.0-35b-uncensored-nvfp4-dflash-vllm