


Sparkrun intelligence scores, failures as zero

Benchmark

official-qwen3-coder-next-int4-autoround-vllm	0.84	0.92	0.92	0.61	0.64	0.53	0.73	0.67	0.76	0.73
eugr-qwen3-coder-next-int4-autoround	0.84	0.92	0.92	0.60	0.64	0.53	0.70	0.66	0.76	0.73
nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	0.84	0.93	0.93	0.61	0.64	0.54	0.72	0.46	0.78	0.72
official-qwen3.6-27b-fp8-vllm	0.86	0.60	0.61	0.60	0.60	0.61	0.78	0.85	0.61	0.68
official-qwen3.6-35b-a3b-fp8-vllm	0.85	0.35	0.36	0.55	0.54	0.59	0.78	0.57	0.69	0.59
eugr-nemotron-3-nano-nvfp4	0.71	0.83	0.46	0.46	0.48	0.53	0.65	0.23	0.59	0.55
aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	0.84	0.12	0.14	0.53	0.53	0.59	0.72	0.62	0.58	0.52
aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	0.79	0.07	0.08	0.50	0.52	0.56	0.71	0.76	0.67	0.52
qwen-qwen3.6-35b-a3b-fp8-mtp-vllm	0.75	0.35	0.37	0.30	0.33	0.33	0.43	0.60	0.68	0.46
official-qwen3.6-35b-a3b-fp8-mtp-vllm	0.32	0.38	0.40	0.21	0.25	0.23	0.32	0.53	0.50	0.35
official-qwen3.6-27b-fp8-mtp-vllm	0.00	0.59	0.60	0.00	0.00	0.25	0.29	0.77	0.40	0.32
official-qwen3-vl-reranker-8b-vllm	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
unsloth-deepseek-v4-flash-671b-gguf-udiq2m-llama-cpp	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
qwen-qwen3.6-27b-fp8-mtp-atlas	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
nvidia-qwen3.6-27b-nvfp4-mtp-vllm	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
official-qwen3-vl-embedding-8b-vllm	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-qwen3.5-397b-int4-autoround	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-openai-gpt-oss-120b	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-glm-4.7-flash-awq	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-diffusion-gemma-nvfp4	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-diffusion-gemma-nvfp4-thinking	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
eugr-diffusion-gemma-bf16	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
unsloth-qwen3.6-27b-nvfp4-mtp-atlas	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00




Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3-coder-next-int4-autoround	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-27b-fp8-vllm
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm
ARC acc	science reasoning	official-qwen3-coder-next-int4-autoround-vllm	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm
ARC norm	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3-coder-next-int4-autoround
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.6-35b-a3b-fp8-vllm	qwen-qwen3.6-35b-a3b-fp8-mtp-vllm


Ranking rule:

- Failures counted as 0
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order

Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size



Sparkrun benchmark TPS by prompt size

