


Sparkrun intelligence scores (grey = test failed or could not run)

Benchmark




Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	atlas-qwen3.6-35b-a3b-fp8-nvfp4	atlas-qwen3.6-35b-a3b-fp8-mtp
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	atlas-qwen3.6-35b-a3b-fp8-mtp	atlas-qwen3.6-35b-a3b-fp8-nvfp4	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC acc	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC norm	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	eugr-nemotron-3-nano-nvfp4


Ranking rule:

- Failures counted as 0 (shown grey with a code)
- FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order
- Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



The chart displays the performance of various models across different prompt token lengths. The x-axis represents the number of prompt tokens, ranging from 256B down to 254K. The y-axis represents a metric, ranging from 0 to 350. The chart shows that for many models, the metric increases significantly as the number of prompt tokens decreases, particularly for tokens below 16K. Some models show a sharp increase in the metric as the number of tokens decreases below 16K, while others remain relatively stable or show a more gradual increase.

Prompt tokens	Model 1 (Dark Blue)	Model 2 (Orange)	Model 3 (Light Blue)	Model 4 (Dark Blue)	Model 5 (Purple)	Model 6 (Green)	Model 7 (Dark Green)	Model 8 (Light Green)	Model 9 (Yellow)	Model 10 (Pink)	Model 11 (Dark Green)	Model 12 (Light Green)	Model 13 (Yellow)	Model 14 (Pink)	Model 15 (Dark Green)
256B	155	25	80	55	45	30	25	20	15	10	5	5	5	5	5
1K	75	25	75	55	55	30	25	20	15	10	5	5	5	5	5
4K	80	30	85	55	45	40	25	20	15	10	5	5	5	5	5
8K	85	55	90	55	75	55	25	20	15	10	5	5	5	5	5
16K	110	80	95	55	110	30	25	20	15	10	5	5	5	5	5
32K	180	120	120	55	180	35	25	20	15	10	5	5	5	5	5
64K	240	155	65	55	240	35	25	20	15	10	5	5	5	5	5
127K	305	250	75	55	305	40	30	20	15	10	5	5	5	5	5
254K	340	340	75	55	95	45	35	25	15	10	5	5	5	5	5

Prompt tokens