





Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
GSM8K strict	math reasoning	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	eugr-nemotron-3-nano-nvfp4	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm	official-qwen3.8-27b-fp8-mtp-vllm
GSM8K flex	math reasoning	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm	official-qwen3.8-27b-fp8-mtp-vllm	official-qwen3.6-27b-fp8-vllm
ARC acc	science reasoning	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie
ARC norm	science reasoning	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm
HellaSwag acc	commonsense	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm
HumanEval	coding	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	official-qwen3.6-27b-fp8-vllm	0xBakeer-qwen3.8-27b-nvfp4-dspark-vllm
MBPP	coding	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-mtp-vllm	official-qwen3.6-27b-fp8-mtp-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	official-qwen3.6-35b-a3b-fp8-mtp-vllm	official-qwen3.6-35b-a3b-fp8-vllm


Ranking rule:

- Failures counted as 0 (shown grey with a code)
- FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order
- Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark TPOT by prompt size

