


Sparkrun intelligence scores (grey = test failed or could not run)

Benchmark




Top 5 models per test


| Test name      | Purpose           | 1st                                              | 2nd                                              | 3rd                                                              | 4th                                                              | 5th                                              |
|----------------|-------------------|--------------------------------------------------|--------------------------------------------------|------------------------------------------------------------------|------------------------------------------------------------------|--------------------------------------------------|
| MMLU           | general knowledge | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae | official-qwen3.6-27b-fp8-vllm                    | community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie | official-qwen3.6-35b-a3b-fp8-vllm                                | eugr-qwen3-coder-next-int4-autoround             |
| GSM8K strict   | math reasoning    | nvidia-nemotron-3-puzzle-75b-a9b                 | eugr-qwen3-coder-next-int4-autoround             | community-north-mini-code-1.0-fp8-vllm-XanuNetworks              | community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks            | poolside-Laguna-S-2-1-NVFP4                      |
| GSM8K flex     | math reasoning    | nvidia-nemotron-3-puzzle-75b-a9b                 | eugr-qwen3-coder-next-int4-autoround             | community-north-mini-code-1.0-fp8-vllm-XanuNetworks              | community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks            | eugr-nemotron-3.5-lightning                      |
| ARC acc        | science reasoning | eugr-nemotron-3.5-lightning                      | nvidia-nemotron-3-puzzle-75b-a9b                 | eugr-qwen3-coder-next-int4-autoround                             | official-qwen3.6-27b-fp8-vllm                                    | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae |
| ARC norm       | science reasoning | eugr-nemotron-3.5-lightning                      | nvidia-nemotron-3-puzzle-75b-a9b                 | eugr-qwen3-coder-next-int4-autoround                             | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae                 | official-qwen3.6-27b-fp8-vllm                    |
| HellaSwag acc  | commonsense       | official-qwen3.6-27b-fp8-vllm                    | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae | official-qwen3.6-35b-a3b-fp8-vllm                                | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui                    | aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm |
| HellaSwag norm | commonsense       | official-qwen3.6-27b-fp8-vllm                    | aeon-7-qwen3.6-35b-a3b-hnvfp4-dflash-vllm        | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae                 | official-qwen3.6-35b-a3b-fp8-vllm                                | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui    |
| HumanEval      | coding            | official-qwen3.6-27b-fp8-vllm                    | aeon-7-qwen3.6-27b-uncenvfp4-dflash-vllm         | eugr-qwen3-coder-next-int4-autoround                             | official-qwen3.6-27b-fp8-mtp-vllm                                | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae |
| MBPP           | coding            | nvidia-nemotron-3-puzzle-75b-a9b                 | eugr-qwen3-coder-next-int4-autoround             | poolside-Laguna-S-2-1-NVFP4                                      | community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie | eugr-nemotron-3.5-lightning                      |
| BFCL overall   | tool calling      | official-qwen3.6-27b-fp8-vllm                    | official-qwen3.6-27b-fp8-mtp-vllm                | eugr-nemotron-3.5-lightning                                      | community-north-mini-code-1.0-fp8-vllm-XanuNetworks              | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui    |

- Ranking rule:
- Failures counted as 0 (shown grey with a code)
  - FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
  - Top 5 chosen by score for each test
  - Ties broken by Mean, then current order
  - Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)





The chart displays the relationship between the number of prompt tokens and a specific metric for various models. The x-axis, labeled 'Prompt tokens', includes categories: 256B, 1K, 4K, 16K, 64K, 127K, and 254K. The y-axis ranges from 0 to 350. The data is represented by multiple colored lines, each corresponding to a different model. Most models show a relatively flat trend until 16K tokens, after which they exhibit a sharp increase. The orange line shows the most dramatic increase, reaching a value of approximately 340 at 254K tokens. The light green line also shows a significant increase, reaching about 305 at 127K tokens. Other models show more moderate increases or remain relatively flat.

| Prompt tokens | Orange | Light Green | Purple | Dark Blue | Yellow | Dark Green | Light Blue | Pink | Grey | Dark Purple | Red | Teal |
|---------------|--------|-------------|--------|-----------|--------|------------|------------|------|------|-------------|-----|------|
| 256B          | 25     | 80          | 130    | 40        | 45     | 55         | 30         | 15   | 20   | 20          | 15  | 30   |
| 1K            | 25     | 65          | 130    | 35        | 50     | 55         | 25         | 15   | 20   | 20          | 15  | 30   |
| 4K            | 30     | 65          | 130    | 45        | 50     | 55         | 25         | 15   | 20   | 20          | 15  | 30   |
| 16K           | 85     | 110         | 130    | 30        | 50     | 55         | 25         | 15   | 20   | 20          | 15  | 30   |
| 64K           | 155    | 245         | 140    | 45        | 55     | 55         | 30         | 15   | 20   | 20          | 15  | 30   |
| 127K          | 250    | 305         | 155    | 75        | 60     | 55         | 35         | 15   | 20   | 20          | 15  | 30   |
| 254K          | 340    | 85          | 165    | 95        | 65     | 55         | 40         | 15   | 20   | 20          | 15  | 30   |

Prompt tokens