


Sparkrun intelligence scores, failures as zero

Benchmark




Score

Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-35b-a3b-fp8-mtp-vllm
GSM8K strict	math reasoning	eugr-qwen3-coder-next-int4-autoround	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-27b-fp8-mtp-vllm	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-mtp-vllm
GSM8K flex	math reasoning	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-35b-a3b-fp8-mtp-vllm
ARC acc	science reasoning	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-27b-fp8-mtp-vllm
ARC norm	science reasoning	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-27b-fp8-mtp-vllm
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-27b-fp8-mtp-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-35b-a3b-fp8-mtp-vllm
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	official-qwen3.6-35b-a3b-fp8-vllm	official-qwen3.6-35b-a3b-fp8-mtp-vllm
MBPP	coding	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-35b-a3b-fp8-vllm	official-qwen3.6-27b-fp8-vllm	eugr-nemotron-3-nano-nvfp4	official-qwen3.6-35b-a3b-fp8-mtp-vllm

Ranking rule:
- Failures counted as 0
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order

Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size



Sparkrun benchmark TPS by prompt size

