


Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
GSM8K strict	math reasoning	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
GSM8K flex	math reasoning	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
ARC acc	science reasoning	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
ARC norm	science reasoning	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
HellaSwag acc	commonsense	official-qwen3.8-27b-fp8-vllm	eugr-deepseek-v4-flash-0731	official-qwen3-coder-next-int4-autoround-vllm		
HellaSwag norm	commonsense	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
HumanEval	coding	official-qwen3.8-27b-fp8-vllm	eugr-deepseek-v4-flash-0731	official-qwen3-coder-next-int4-autoround-vllm		
MBPP	coding	eugr-deepseek-v4-flash-0731	official-qwen3.8-27b-fp8-vllm	official-qwen3-coder-next-int4-autoround-vllm		
BFCL 7-18 of 22	tool calling	official-qwen3-coder-next-int4-autoround-vllm	official-qwen3.8-27b-fp8-vllm	eugr-deepseek-v4-flash-0731		


Ranking rule:

- Failures counted as 0 (shown grey with a code)
- FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order
- Mean covers lm-eval metrics only (excludes BFCL)


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size

