


Sparkrun intelligence scores, failures as zero

Benchmark




Score

Top 5 models per test


Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	eugr-qwen3-coder-next-int4-autoround
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	atlas-qwen3.6-35b-a3b-fp8-nvfp4	atlas-qwen3.6-35b-a3b-fp8-mtp
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	atlas-qwen3.6-35b-a3b-fp8-mtp	atlas-qwen3.6-35b-a3b-fp8-nvfp4	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC acc	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
ARC norm	science reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks
HellaSwag acc	commonsense	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-h-nvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncen-nvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4	eugr-qwen3-coder-next-int4-autoround	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm

Ranking rule:
- Failures counted as 0
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size

