



Sparkrun intelligence scores, failures as zero

Benchmark




Top 5 models per test


| Test name      | Purpose           | 1st                                              | 2nd                                              | 3rd                                                              | 4th                                              | 5th                                                 |
|----------------|-------------------|--------------------------------------------------|--------------------------------------------------|------------------------------------------------------------------|--------------------------------------------------|-----------------------------------------------------|
| MMLU           | general knowledge | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae | official-qwen3.6-27b-fp8-vllm                    | community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie | official-qwen3.6-35b-a3b-fp8-vllm                | eugr-qwen3-coder-next-int4-autoround                |
| GSM8K strict   | math reasoning    | nvidia-nemotron-3-puzzle-75b-a9b-nvfp4           | eugr-qwen3-coder-next-int4-autoround             | community-north-mini-code-1.0-fp8-vllm-XanuNetworks              | atlas-qwen3.6-35b-a3b-fp8-nvfp4                  | atlas-qwen3.6-35b-a3b-fp8-mtp                       |
| GSM8K flex     | math reasoning    | nvidia-nemotron-3-puzzle-75b-a9b-nvfp4           | eugr-qwen3-coder-next-int4-autoround             | atlas-qwen3.6-35b-a3b-fp8-mtp                                    | atlas-qwen3.6-35b-a3b-fp8-nvfp4                  | community-north-mini-code-1.0-fp8-vllm-XanuNetworks |
| ARC acc        | science reasoning | nvidia-nemotron-3-puzzle-75b-a9b-nvfp4           | eugr-qwen3-coder-next-int4-autoround             | official-qwen3.6-27b-fp8-vllm                                    | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae | community-north-mini-code-1.0-fp8-vllm-XanuNetworks |
| ARC norm       | science reasoning | nvidia-nemotron-3-puzzle-75b-a9b-nvfp4           | eugr-qwen3-coder-next-int4-autoround             | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae                 | official-qwen3.6-27b-fp8-vllm                    | community-north-mini-code-1.0-fp8-vllm-XanuNetworks |
| HellaSwag acc  | commonsense       | official-qwen3.6-27b-fp8-vllm                    | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae | official-qwen3.6-35b-a3b-fp8-vllm                                | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui    | aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm    |
| HellaSwag norm | commonsense       | official-qwen3.6-27b-fp8-vllm                    | aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae                 | official-qwen3.6-35b-a3b-fp8-vllm                | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui       |
| HumanEval      | coding            | official-qwen3.6-27b-fp8-vllm                    | aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm  | eugr-qwen3-coder-next-int4-autoround                             | official-qwen3.6-27b-fp8-mtp-vllm                | community-qwen3.5-27b-fp8-dflash-vllm-banana-bae    |
| MBPP           | coding            | nvidia-nemotron-3-puzzle-75b-a9b-nvfp4           | eugr-qwen3-coder-next-int4-autoround             | community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie | official-qwen3.6-35b-a3b-fp8-vllm                | aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm     |
| BFCL overall   | tool calling      | official-qwen3.6-27b-fp8-vllm                    | official-qwen3.6-27b-fp8-mtp-vllm                | community-north-mini-code-1.0-fp8-vllm-XanuNetworks              | community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui    | eugr-nemotron-3-nano-nvfp4                          |

- Ranking rule:
- Failures counted as 0
  - Top 5 chosen by score for each test
  - Ties broken by Mean, then current order
  - Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)





### Sparkrun benchmark TPOT by prompt size

