


Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
GSM8K strict	math reasoning	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
GSM8K flex	math reasoning	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
ARC acc	science reasoning	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-deepseek-v4-flash-0731-b1-vllm	eugr-deepseek-v4-flash-0731	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
ARC norm	science reasoning	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
HellaSwag acc	commonsense	eugr-qwen3.8-flash-next-nvfp4-cluster	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm
HellaSwag norm	commonsense	eugr-qwen3.8-flash-next-nvfp4-cluster	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm
HumanEval	coding	eugr-qwen3.8-flash-next-nvfp4-cluster	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm
MBPP	coding	eugr-deepseek-v4-flash-0731	official-deepseek-v4-flash-0731-b1-vllm	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm	eugr-qwen3.8-flash-next-nvfp4-cluster
BFCL 17 of 22	tool calling	official-deepseek-v4-flash-0731-b1-vllm	eugr-deepseek-v4-flash-0731	eugr-qwen3.8-flash-next-nvfp4-cluster	miaai-glm-5.3-flash-exl3-4bpw-dflash2-k7-2x-vllm	official-qwen3-coder-next-int4-autoround-vllm


Ranking rule:

- Failures counted as 0 (shown grey with a code)
- FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order
- Mean covers lm-eval metrics only (excludes BFCL)


Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



Sparkrun benchmark time to first token by prompt size



Sparkrun benchmark TPOT by prompt size

