





Top 5 models per test

Test name	Purpose	1st	2nd	3rd	4th	5th
MMLU	general knowledge	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	official-qwen3.6-35b-a3b-fp8-vllm	qwen-qwen3.8-27b-fp8-vllm
GSM8K strict	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	poolside-Laguna-S-2-1-NVFP4
GSM8K flex	math reasoning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	community-north-mini-code-1.0-fp8-vllm-XanuNetworks	community-north-mini-code-1.0-nvfp4-vllm-XanuNetworks	eugr-nemotron-3.5-lightning
ARC acc	science reasoning	eugr-nemotron-3.5-lightning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
ARC norm	science reasoning	eugr-nemotron-3.5-lightning	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-27b-fp8-vllm
HellaSwag acc	commonsense	qwen-qwen3.8-27b-fp8-vllm	official-qwen3.6-27b-fp8-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	community-qwen3.6-35b-a3b-fp8-vllm-povedaaqui
HellaSwag norm	commonsense	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-35b-a3b-heretic-nvfp4-dflash-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae	official-qwen3.6-35b-a3b-fp8-vllm	qwen-qwen3.8-27b-fp8-vllm
HumanEval	coding	official-qwen3.6-27b-fp8-vllm	aeon-7-qwen3.6-27b-uncensored-nvfp4-dflash-vllm	eugr-qwen3-coder-next-int4-autoround	official-qwen3.6-27b-fp8-mtp-vllm	community-qwen3.5-27b-fp8-dflash-vllm-banana-bae
MBPP	coding	nvidia-nemotron-3-puzzle-75b-a9b-nvfp4-vllm	eugr-qwen3-coder-next-int4-autoround	poolside-Laguna-S-2-1-NVFP4	community-qwen3.6-35b-a3b-autoround-int4-dflash-vllm-cipherfoxie	eugr-nemotron-3.5-lightning
BFCL overall	tool calling	official-qwen3.6-27b-fp8-vllm	official-qwen3.6-27b-fp8-mtp-vllm	eugr-nemotron-3.5-lightning	qwen-qwen3.8-27b-fp8-vllm	community-north-mini-code-1.0-fp8-vllm-XanuNetworks


Ranking rule:

- Failures counted as 0 (shown grey with a code)
- FAILED: server never started; UNSUP: server lacks echo+logprobs for multiple-choice tasks; OOM/CRASH/HANG/ERROR: from the report
- Top 5 chosen by score for each test
- Ties broken by Mean, then current order
- Mean covers lm-eval metrics only (excludes BFCL)

Sparkrun benchmark TPS by prompt size



Sparkrun benchmark aggregate TPS by concurrency (1024-token prompts)



[illegible]

Prompt tokens